


# Machine Learning

## CS7052

### Lecture 7, Decision Trees

Dr. Elaheh Homayounvala

Week 7



# Outline of today's lecture

- Review week 5
- Supervised learning, Decision Trees

# Linear Models (and Gradient Descent)

# What we covered so far

- Supervised learning

1. kNN
2. Linear Models

kNN and linear models can be applied in classification as well as regression


- Unsupervised learning

- Clustering

# k-Nearest Neighbours Algorithm


Classification:

- 1 Find  $k$  closest objects to the predicted object  $x$  in the training set.
- 2 Associate  $x$  the most frequent class among its  $k$  neighbours.



Regression:

- 1 Find  $k$  closest objects to the predicted object  $x$  in the training set.
- 2 Associate  $x$  average output of its  $k$  neighbours.



# Linear Models

- Linear models:

$$\hat{y} = w[0] * x[0] + w[1] * x[1] + ... + w[p] * x[p] + b$$

- Simple form of cost or loss function:

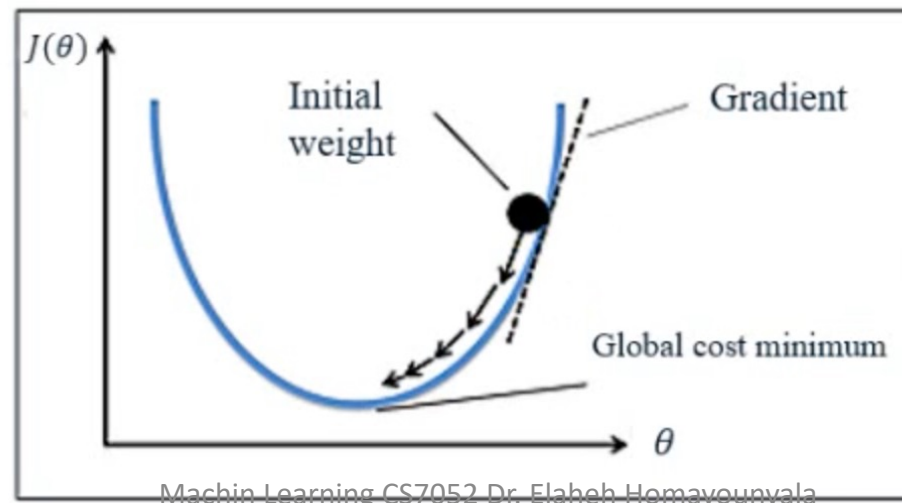
$$\sum_{i=1}^M (y_i - \hat{y}_i)^2 = \sum_{i=1}^M \left( y_i - \sum_{j=0}^p w_j \times x_{ij} \right)^2$$

Ryan Holiday

The data-set has M instances and p features

# Linear Models, Gradient Descent

- We look for optimising  $w$  and  $b$  such that it minimises the cost function
- Gradient Descent is one of the ways for cost minimisation
- What are the others?



# Practical skills in ML

Jupyter Notebook

numpy, pandas, matplotlib, scikitlearn libraries



# Practical ML

- Choose a dataset and a model and a tool
- Construct/Learn a model based on training data
- Use the model to make predictions for test data
- Evaluate the model (accuracy)
- Tune the model or choose another if necessary

# What is a decision tree?

## Information Gain, Gain Ratio and Gini index

## Parameters, Strengths and Weaknesses

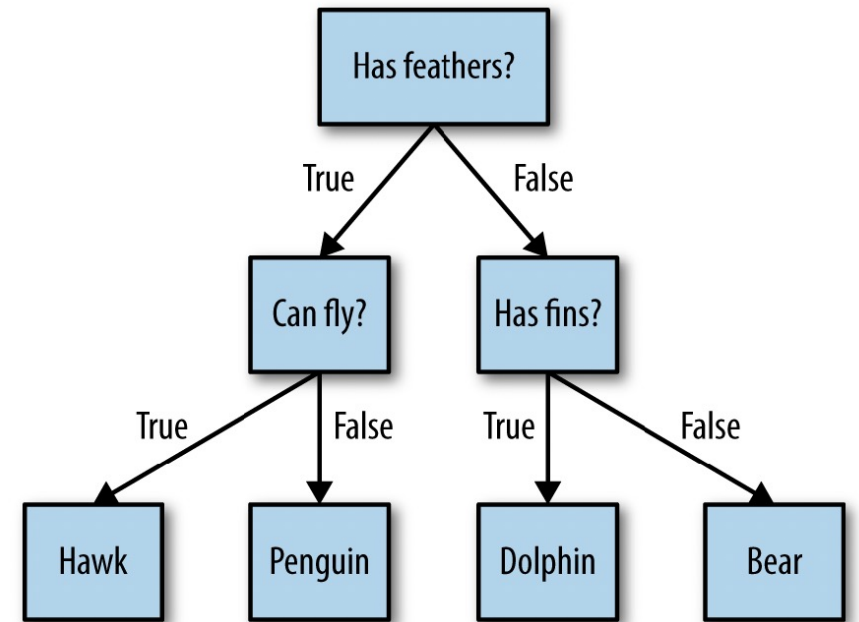
# Decision Trees

- Widely used models for classification and regression tasks
- They learn a hierarchy of if/else questions, leading to a decision

# Decision trees

Distinguish between the following four animals

- Bears
  - Hawks
  - Penguins
  - dolphins.
- 
- Muller and Guido's book, page 73
  - Animal pictures from pixabay.com



*Figure 2-22. A decision tree to distinguish among several animals*



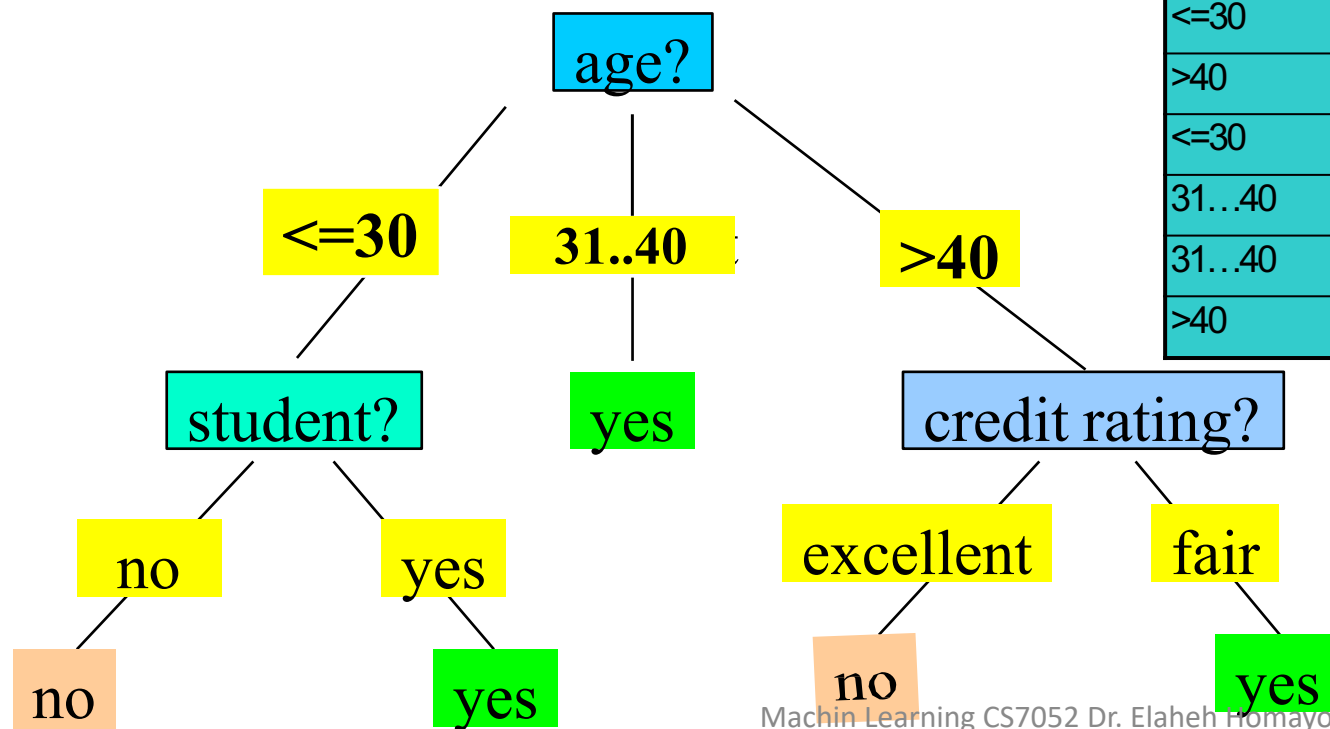
# Nodes and Edges in Decision Trees

- Each node in the tree either represents
  - a question or
  - a terminal node (also called a leaf) that contains the answer
- The edges connect the answers to a question with the next question you would ask.
- It is an upside-down tree, root at top!



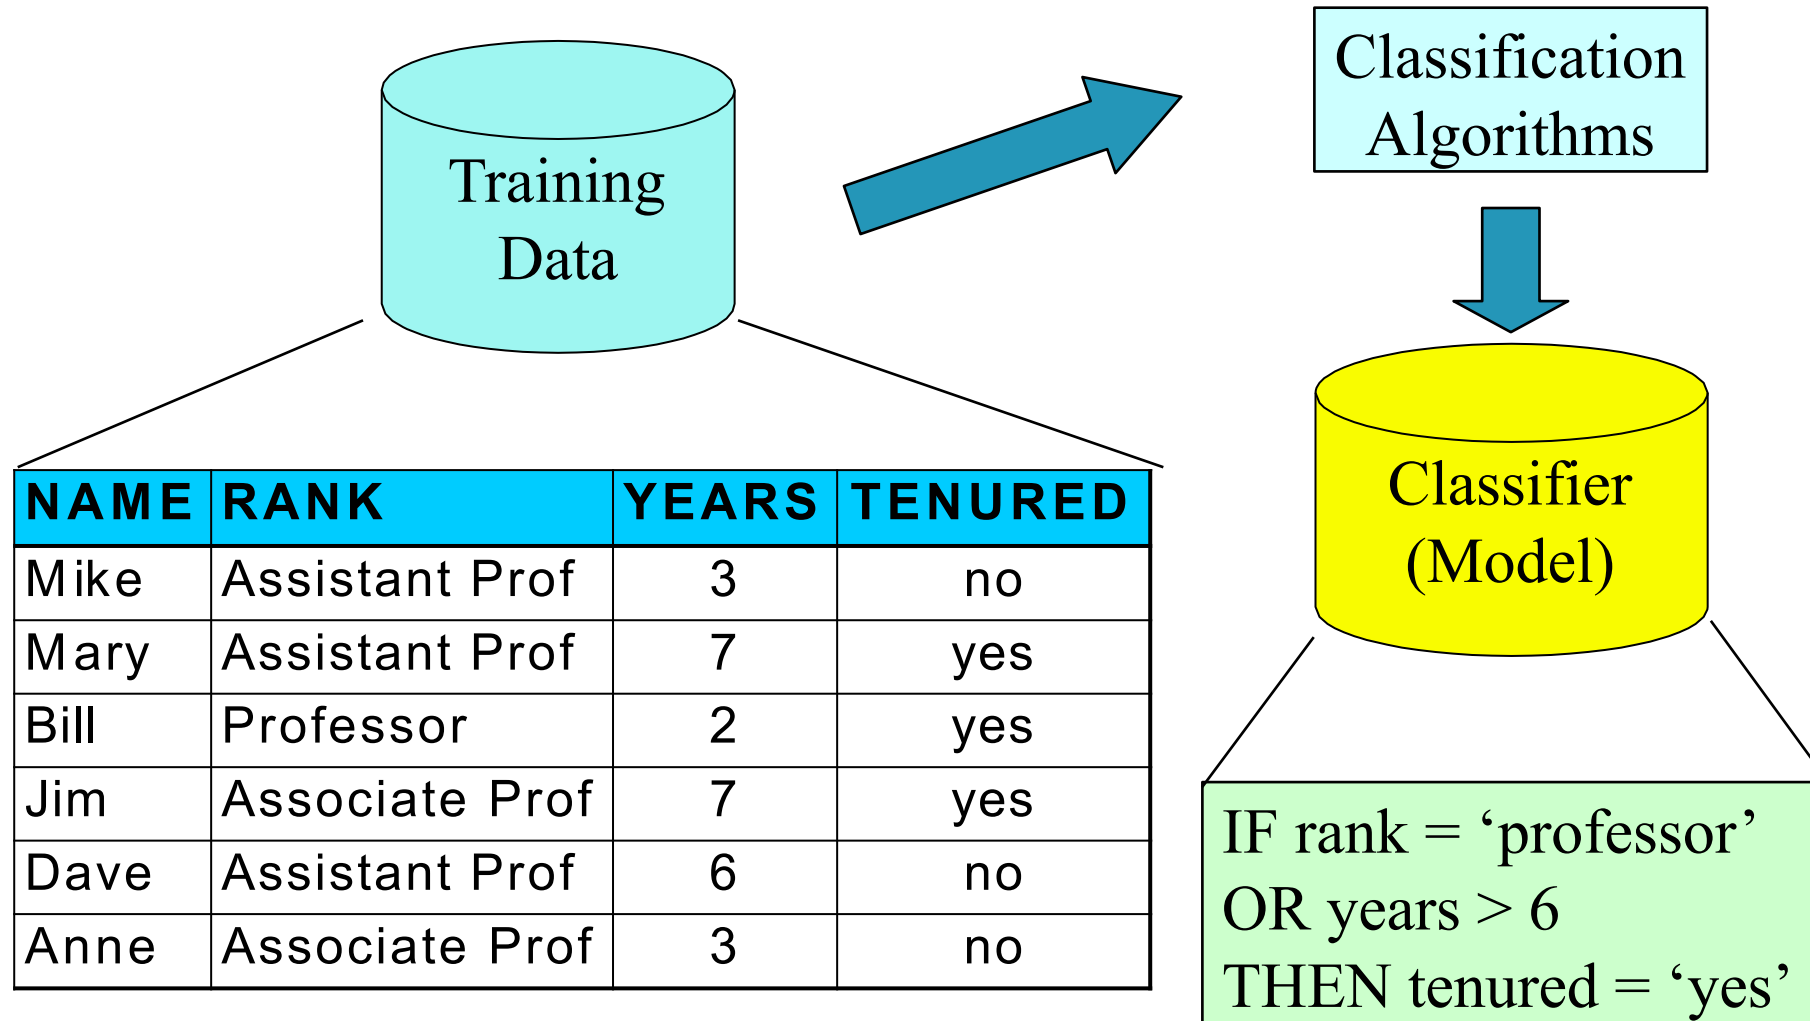
# Decision Tree, An Example

- Training dataset: Buys\_computer
- Decision Tree:

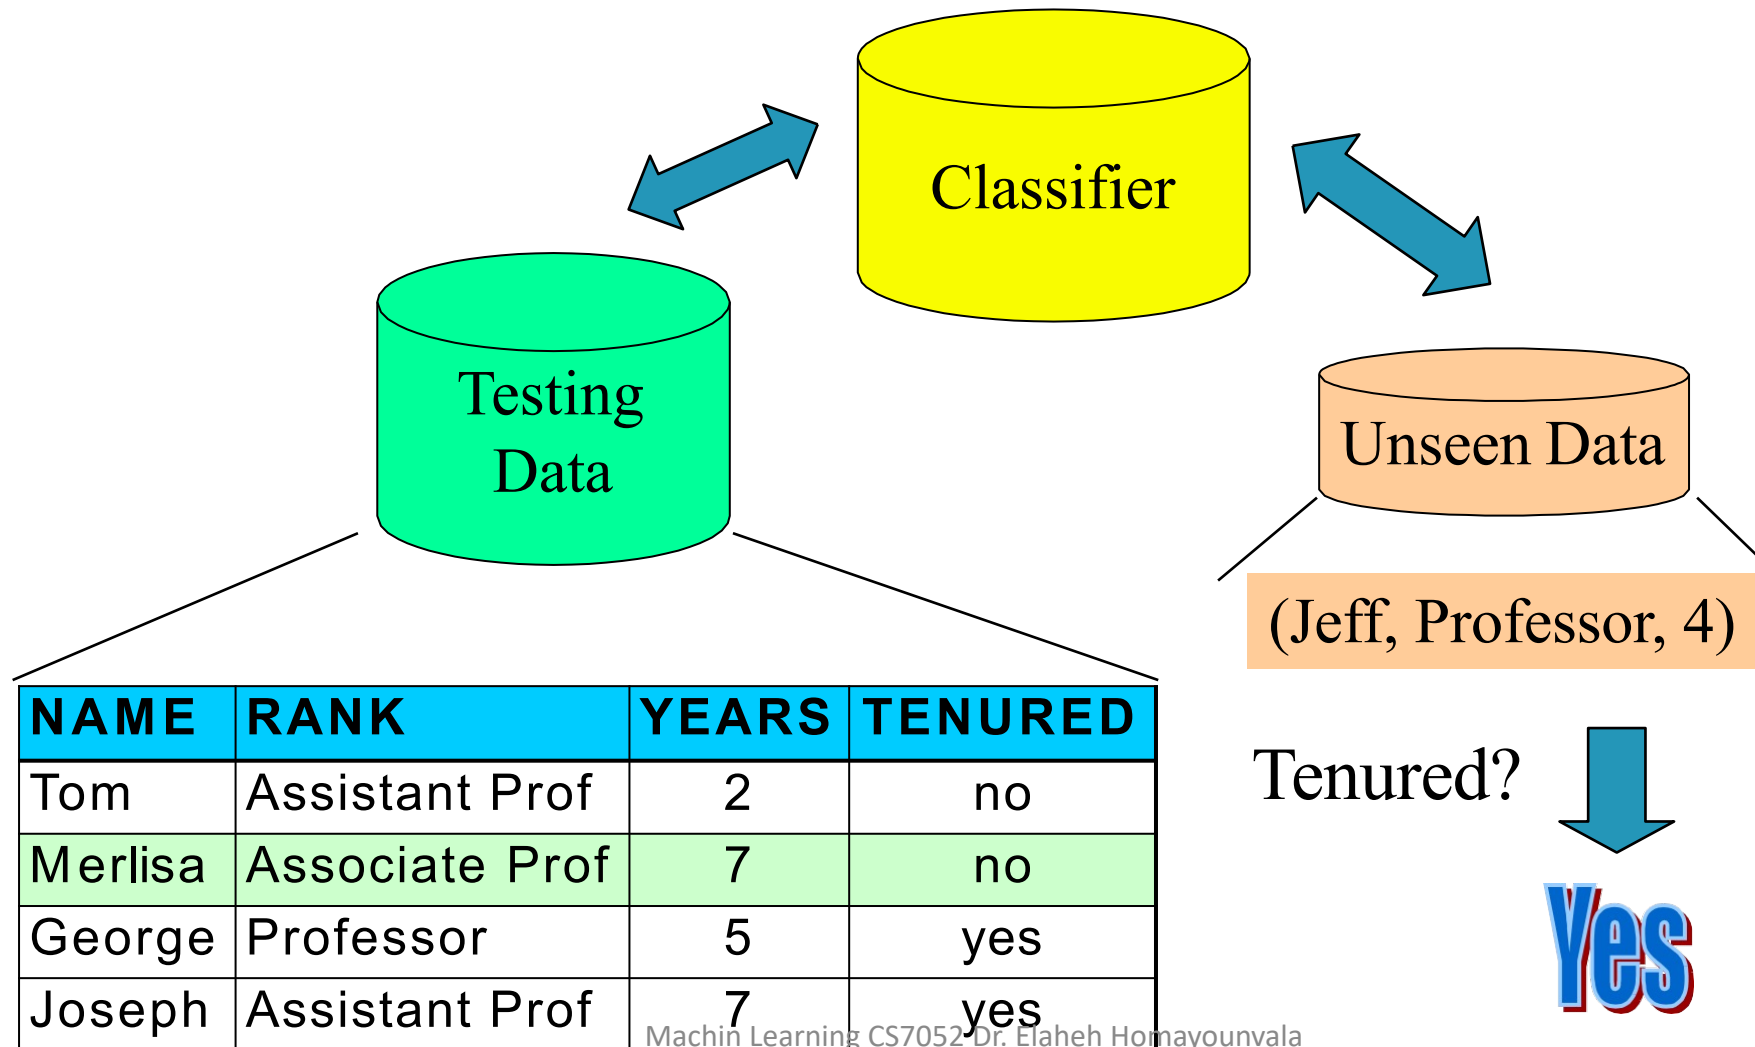


| age    | income | student | credit_rating | buys_computer |
|--------|--------|---------|---------------|---------------|
| <=30   | high   | no      | fair          | no            |
| <=30   | high   | no      | excellent     | no            |
| 31..40 | high   | no      | fair          | yes           |
| >40    | medium | no      | fair          | yes           |
| >40    | low    | yes     | fair          | yes           |
| >40    | low    | yes     | excellent     | no            |
| 31..40 | low    | yes     | excellent     | yes           |
| <=30   | medium | no      | fair          | no            |
| <=30   | low    | yes     | fair          | yes           |
| >40    | medium | yes     | fair          | yes           |
| <=30   | medium | yes     | excellent     | yes           |
| 31..40 | medium | no      | excellent     | yes           |
| 31..40 | high   | yes     | fair          | yes           |
| >40    | medium | no      | excellent     | no            |

# Process 1, model construction (induction)



# Process 2, using the model in prediction





# Building/Learning a Decision Tree


- Learning a decision tree means learning the sequence of if/else questions that gets us to the true answer most quickly.
- In ML if/else questions are called “tests”.

# Learning a Decision Tree

- To build a tree, the algorithm searches over all possible tests and finds the one that is most informative about the target variable.
- Continue and choose the next test recursively.
- The recursive partitioning of the data is repeated until
  - Each region in the partition (each leaf) only contains a single target value.
- A leaf of the tree that contains data points that all share the same target value is called pure.

# Building a Decision Tree, example

- Muller and Guido's book, page 74



*Figure 2-23. Two-moons dataset on which the decision tree will be built*

# Decision Boundary of Tree, depth 1

- Muller and Guido's book, page 74





Figure 2-24. Decision boundary of tree with depth 1 (left) and corresponding tree (right)

# Decision Boundary of Tree, depth 2

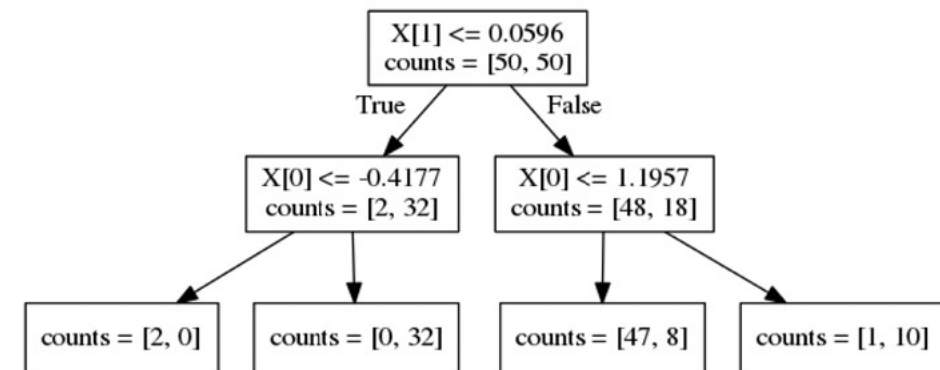



Figure 2-25. Decision boundary of tree with depth 2 (left) and corresponding decision tree (right)

# Algorithm for Decision Tree Induction, more detailed

- Basic algorithm (a greedy algorithm)
  - Tree is constructed in a top-down **recursive divide-and-conquer** manner
  - At start, all the training examples are at the root
  - Attributes are categorical (if continuous-valued, they are discretised in advance)
  - Examples are partitioned recursively based on selected attributes
  - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)

# Algorithm for Decision Tree Induction, cont.

## Conditions for stopping partitioning

- All samples for a given node belong to the same class
- There are no remaining attributes for further partitioning **majority voting** is employed for classifying the leaf
- There are no samples left


# How do we choose the order of test attributes?

- To build a tree, the algorithm searches over all possible tests and finds the one that is most informative about the target variable.
- How can we identify the most informative test?
- Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)
- What are those heuristic or statistical measures?
- What is information gain?
- First let's learn a new concept: Entropy




# Entropy

- A measure of uncertainty for a random variable
  - Calculation: For a discrete random variable  $Y$  taking  $m$  distinct values  $\{y_1, \dots, y_m\}$ ,
    - $H(Y) = -\sum_{i=1}^m p_i \log(p_i)$ , where  $p_i = P(Y = y_i)$




Bucket 1

Entropy = 0



Bucket 2

Entropy = 0.81125




Bucket 3

Entropy = 1

# Brief review of Entropy

- Entropy (Information Theory)
  - A measure of uncertainty associated with a random variable
  - Calculation: For a discrete random variable  $Y$  taking  $m$  distinct values  $\{y_1, \dots, y_m\}$ ,
    - $H(Y) = -\sum_{i=1}^m p_i \log(p_i)$ , where  $p_i = P(Y = y_i)$
  - Interpretation:
    - Higher entropy => higher uncertainty
    - Lower entropy => lower uncertainty
- Conditional Entropy
  - $H(Y|X) = \sum_x p(x)H(Y|X = x)$

**m = 2**

# Attribute selection measure

## Information Gain (ID3/C4.5)

- Select the attribute with the highest information gain
- Let  $p_i$  be the probability that an arbitrary tuple in  $D$  belongs to class  $C_i$ , estimated by  $|C_{i,D}| / |D|$
- Missing information (entropy) needed to classify a tuple in  $D$ :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$

- Missing Information after using  $A$  to split  $D$  into  $v$  partitions :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

- Information gained by branching on attribute  $A$

$$Gain(A) = Info(D) - Info_A(D)$$

# Attribute selection, Information Gain

- Class P: buys\_computer = "yes"
- Class N: buys\_computer = "no"

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

| age     | p <sub>i</sub> | n <sub>i</sub> | I(p <sub>i</sub> , n <sub>i</sub> ) |
|---------|----------------|----------------|-------------------------------------|
| <=30    | 2              | 3              | 0.971                               |
| 31...40 | 4              | 0              | 0                                   |
| >40     | 3              | 2              | 0.971                               |

| age     | income | student | credit_rating | buys_computer |
|---------|--------|---------|---------------|---------------|
| <=30    | high   | no      | fair          | no            |
| <=30    | high   | no      | excellent     | no            |
| 31...40 | high   | no      | fair          | yes           |
| >40     | medium | no      | fair          | yes           |
| >40     | low    | yes     | fair          | yes           |
| >40     | low    | yes     | excellent     | no            |
| 31...40 | low    | yes     | excellent     | yes           |
| <=30    | medium | no      | fair          | no            |
| <=30    | low    | yes     | fair          | yes           |
| >40     | medium | yes     | fair          | yes           |
| <=30    | medium | yes     | excellent     | yes           |
| 31...40 | medium | no      | excellent     | yes           |
| 31...40 | high   | yes     | fair          | yes           |
| >40     | medium | no      | excellent     | no            |

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0)$$

$$+ \frac{5}{14} I(3,2) = 0.694$$

$\frac{5}{14} I(2,3)$  means "age <=30" has 5 out of 14 samples, with 2 yes'es and 3 no's. Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit\_rating) = 0.048$$

# Computing Information Gain for Continuous-valued attributes

- Let attribute  $A$  be a continuous-valued attribute
- Must determine the best split point for  $A$ 
  - Sort the value  $A$  in increasing order
  - Typically, the midpoint between each pair of adjacent values is considered as a possible *split point*
    - $(a_i + a_{i+1})/2$  is the midpoint between the values of  $a_i$  and  $a_{i+1}$
  - The point with the *minimum expected information requirement* for  $A$  is selected as the split-point for  $A$
- Split:
  - $D_1$  is the set of tuples in  $D$  satisfying  $A \leq \text{split-point}$ , and  $D_2$  is the set of tuples in  $D$  satisfying  $A > \text{split-point}$

# Gain ratio for attribute selection

- Information gain measure is biased towards attributes with a large number of values
- C4.5 (a successor of ID3) uses gain ratio to overcome the problem (normalisation to information gain)

$$SplitInfo_A(D) = -\sum_{j=1}^v \frac{|D_j|}{|D|} \log_2 \left( \frac{|D_j|}{|D|} \right)$$

- $GainRatio(A) = Gain(A)/SplitInfo(A)$

- Example  $SplitInfo_{income}(D) = -\frac{4}{14} \times \log_2 \left( \frac{4}{14} \right) - \frac{6}{14} \times \log_2 \left( \frac{6}{14} \right) - \frac{4}{14} \times \log_2 \left( \frac{4}{14} \right) = 1.557$

- $gain\_ratio(income) = 0.029/1.557 = 0.019$

- The attribute with the maximum gain ratio is selected as the splitting attribute

# Gini Index (Cart, IBM Intelligent Miner)

- If a data set  $D$  contains examples from  $n$  classes, gini index,  $gini(D)$  is defined as

$$gini(D) = 1 - \sum_{j=1}^n p_j^2$$

where  $p_j$  is the relative frequency of class  $j$  in  $D$

- If a data set  $D$  is split on  $A$  into two subsets  $D_1$  and  $D_2$ , the gini index  $gini(D)$  is defined as

$$gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$$

- Reduction in Impurity:

$$\Delta gini(A) = gini(D) - gini_A(D)$$

- The attribute provides the smallest  $gini_{split}(D)$  (or the largest reduction in impurity) is chosen to split the node (*need to enumerate all the possible splitting points for each attribute*)

# Computation of Gini index

- Example: D has 9 tuples in `buys_computer = "yes"` and 5 in "no"

$$gini(D) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$

- Suppose the attribute `income` partitions D into 10 in  $D_1: \{\text{low, medium}\}$  and 4 in  $D_2$

$$\begin{aligned} gini_{income \in \{\text{low, medium}\}}(D) &= \left(\frac{10}{14}\right) Gini(D_1) + \left(\frac{4}{14}\right) Gini(D_2) \\ &= \frac{10}{14} \left(1 - \left(\frac{7}{10}\right)^2 - \left(\frac{3}{10}\right)^2\right) + \frac{4}{14} \left(1 - \left(\frac{2}{4}\right)^2 - \left(\frac{2}{4}\right)^2\right) \\ &= 0.443 \\ &= Gini_{income \in \{\text{high}\}}(D). \end{aligned}$$

$Gini_{\{\text{low, high}\}}$  is 0.458;  $Gini_{\{\text{medium, high}\}}$  is 0.450. Thus, split on the  $\{\text{low, medium}\}$  (and  $\{\text{high}\}$ ) since it has the lowest Gini index



# Comparing attribute selection methods

- The three measures in general return good result.
  - **Information gain:**
    - biased towards multivalued attributes
  - **Gain ratio:**
    - tends to prefer unbalanced splits in which one partition is much smaller than the others
  - **Gini index:**
    - biased to multivalued attributes
    - has difficulty when number of classes is large
    - tends to favour tests that result in equal-sized partitions and purity in both partitions

# Decision Trees for Regression

- Similar to decision trees for classification but
- Decision Tree Regressor is not able to extrapolate (make predictions outside of the range of the training data)
- Muller and Guido's book, page 84




Figure 2-32. Comparison of predictions made by a linear model and predictions made by a regression tree on the RAM price data

# Overfitting and Tree Pruning

- Overfitting: An induced tree may overfit the training data
  - Too many branches, some may reflect anomalies due to noise or outliers
  - Poor accuracy for unseen samples
- Two approaches to avoid overfitting
  - Pre-pruning: *Halt tree construction early* - do not split a node if this would result in the goodness measure falling below a threshold
    - Difficult to choose an appropriate threshold
  - Post-pruning: *Remove branches* from a “fully grown” tree— removing or collapsing nodes that contain little information

# Strengths, weaknesses and parameters

# Parameters

- pre-pruning parameters such as:
  - maximum depth of the tree
  - maximum number of leaves
  - a minimum number of points in a node to keep splitting it

# Strengths

- Can be easily visualized and understood by non-experts
- The algorithms are completely invariant to scaling of the data
  - no pre-processing like normalization or standardization of features is needed for decision tree algorithms.
- Work well when you have features that are on completely different scales or a mix of binary and continuous features

# Weaknesses

- They tend to overfit (even with pre-pruning)
- Poor generalisation performance
- Ensemble methods are used instead of single trees

# Ensemble of Decision Trees

Further Reading and possible topics for coursework

- Random Forest
- Gradient boosted regression trees